

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statistics import mode
from scipy.stats import levene # Optional assumption for standard t-
test
from scipy.stats import ttest_ind # To perform t-test

df = pd.read_csv("Apple_Store_Reviews.csv")
df

```

	Review_ID	App_Name	User_Age	Review_Date	Rating	\
0	1	Candy Crush Saga	21	2023-01-16	4	
1	2	Spotify	57	2024-02-01	1	
2	3	TikTok	33	2023-11-30	5	
3	4	Audible	40	2023-04-03	5	
4	5	Spotify	44	2023-05-01	1	
..	
995	996	Headspace	30	2023-11-15	3	
996	997	Duolingo	19	2024-09-27	1	
997	998	Duolingo	38	2023-06-07	5	
998	999	Instagram	52	2024-03-04	4	
999	1000	Audible	25	2024-02-20	2	

	Review_Text	Likes	Device_Type	\
0	Great game, but too many in-game purchases.	70	iPhone 12	
1	Good, but has connection issues sometimes.	49	iPhone SE	
2	Awesome app! Best entertainment content.	98	iPhone 12	
3	Great app, but it's a bit pricey.	74	iPhone 13	
4	Good, but has connection issues sometimes.	47	iPhone SE	
..	
995	Good, but the premium content is expensive.	65	iPhone SE	
996	Disappointing. Hard to follow and buggy.	4	iPhone SE	
997	Excellent for learning new skills!	85	iPhone 11	
998	Great app, but sometimes it lags.	55	iPhone 13	
999	Terrible. Very limited selection of books.	7	iPhone 13	

	Version_Used	Country	Purchase_Amount	Category
0	3.231.19	Australia	0.00	Games
1	4.102.9	Germany	7.15	Music
2	7.52.0	Germany	4.98	Entertainment
3	5.260.15	Australia	0.00	Books
4	4.50.18	Australia	14.31	Music
..
995	6.284.11	US	0.00	Health
996	6.293.8	Canada	7.25	Education
997	10.277.15	Mexico	13.33	Education
998	3.52.20	US	6.37	Social
999	9.150.8	India	8.32	Books

```
[1000 rows x 12 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	Review_ID	1000 non-null	int64
1	App_Name	1000 non-null	object
2	User_Age	1000 non-null	int64
3	Review_Date	1000 non-null	object
4	Rating	1000 non-null	int64
5	Review_Text	1000 non-null	object
6	Likes	1000 non-null	int64
7	Device_Type	1000 non-null	object
8	Version_Used	1000 non-null	object
9	Country	1000 non-null	object
10	Purchase_Amount	1000 non-null	float64
11	Category	1000 non-null	object

```
dtypes: float64(1), int64(4), object(7)
```

```
memory usage: 93.9+ KB
```

```
df.isnull().sum()
```

Review_ID	0
App_Name	0
User_Age	0
Review_Date	0
Rating	0
Review_Text	0
Likes	0
Device_Type	0
Version_Used	0
Country	0
Purchase_Amount	0
Category	0

```
dtype: int64
```

```
df["Review_Date"] = pd.to_datetime(df["Review_Date"])
```

```
df["Rating"].describe()
```

count	1000.000000
mean	2.869000
std	1.467649
min	1.000000
25%	1.000000
50%	3.000000

```

75%          4.000000
max          5.000000
Name: Rating, dtype: float64

mean = df["Rating"].mean()
print(mean)
# Median - Robust to outliers and skewed distributions; provides a
better central tendency for non-symmetrical data.

median = df["Rating"].median()
print(median)

# Mode: Useful when the data has a prominent peak or category that
repeats frequently.

mode = df["Rating"].mode()
print(mode)

2.869
3.0
0    1
Name: Rating, dtype: int64

# Median - Robust to outliers and skewed distributions; provides a
better central tendency for non-symmetrical data.

median = df["Rating"].median()
median

3.0

# Mode: Useful when the data has a prominent peak or category that
repeats frequently.

mode = mode(df["Rating"])
mode


1

df["Rating"].mode()

0    1
Name: Rating, dtype: int64

plt.figure(figsize = (4,4))
sns.histplot(data = df["Rating"], bins = 5, color = "lightblue",
edgecolor = "black", kde = True)
plt.title("Frequency distribution of App Ratings", fontsize = 10)
plt.show()

```



```
# Range = (min_value, max_value)
# IQR = Q3 - Q1(0.75 - 0.25)
'''
```

Conclusion :

The Range highlights the extreme values in the data, while the IQR gives a more focused view of typical purchase behaviors. Based on this, the dataset shows significant variability in purchase amounts, with a notable cluster at lower values.

Insights:

- The data has a large spread, with some high-value purchases pushing the maximum to \$19.97. However, the low Q1 shows a cluster of free or minimal-value purchases.*
- The relatively high IQR (\$10.19) compared to the Range (\$19.97) indicates that the variability in the middle 50% of data is significant.*

This implies diverse spending behaviors among users.

```
'''

minimum = df["Purchase_Amount"].min()
maximum = df["Purchase_Amount"].max()

q1 = df["Purchase_Amount"].quantile(0.25)
q3 = df["Purchase_Amount"].quantile(0.75)

Range = maximum - minimum
```

```

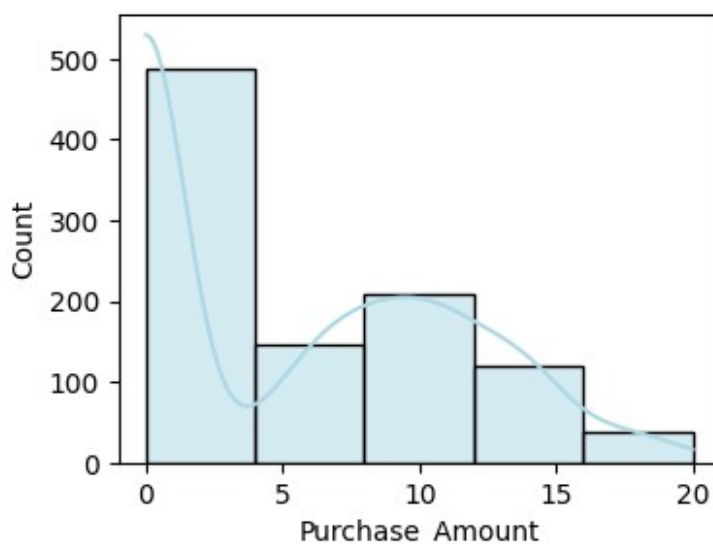
IQR = q3 - q1

print("Maximum Purchase Amount :",maximum)
print("Minimum Purchase Amount : ",minimum)
print("Q1 (25th percentile): ",q1)
print("Q3 (75th percentile):",q3)
print("Range (Max - Min):", Range)
print("Interquartile Range (IQR = Q3 - Q1):", IQR)

plt.figure(figsize = (4,3))
sns.histplot(data = df["Purchase_Amount"], kde = True, bins = 5, color
= "lightblue",edgecolor = "black")
plt.show()

Maximum Purchase Amount : 19.97
Minimum Purchase Amount : 0.0
Q1 (25th percentile): 0.0
Q3 (75th percentile): 10.192499999999999
Range (Max - Min): 19.97
Interquartile Range (IQR = Q3 - Q1): 10.192499999999999

```



```

...
Standard Deviation (28.69):
The standard deviation indicates a moderate level of variability in
the number of likes.
Most reviews tend to receive likes that are within a range of
approximately 28.69 likes above or below the mean.
This suggests that while some reviews perform better or worse than
average, the majority of reviews have a fairly consistent level of
engagement.

Variance (822.85):

```

The variance, being the square of the standard deviation, quantifies the spread of the data but in squared units. The relatively high value confirms the presence of variability, though it is less intuitive to interpret compared to the standard deviation.

Conclusion :

The combination of moderate standard deviation and variance highlights that while there are variations in the number of likes received, most reviews tend to perform within a predictable range around the average.


'''

```
std = df["Likes"].std()
print("Standard deviation : ",std)
var = df["Likes"].var()
print("Variance : ",var)

plt.figure(figsize = (4,4))
sns.histplot(data= df["Likes"], kde = True, bins = 20)
plt.show()
```

Standard deviation : 28.685443672334557

Variance : 822.8546786786787



'''

Interpretation of the Correlation:

Value of Correlation (0.8425):

Correlation values range from -1 to 1:

+1: Perfect positive correlation.

0: No correlation.

-1: Perfect negative correlation.

A correlation of 0.8425 is close to +1, suggesting a strong positive relationship.

What It Means:

As the rating of a review increases, the number of likes it receives also tends to increase.

High ratings are strongly associated with high engagement (likes), indicating that users are more likely to like reviews they perceive as positive or useful.

'''

```
corr = df[["Likes","Rating"]].corr()
```

```
print("Correlation : ", corr)
```


```
plt.figure(figsize = (3,3))
```

```
sns.heatmap(data = corr, annot = True)
```

```
plt.title("Correlation b/w Likes and Ratings", fontsize = 9)
```

```
plt.show()
```

Correlation :	Likes	Rating
Likes	1.000000	0.842541
Rating	0.842541	1.000000



'''

A skewness value close to 0 indicates that the distribution is

approximately symmetrical.

*Skewness Value: The skewness of the app ratings is 0.1018
0.1018, which is close to zero. This indicates an approximately
symmetrical distribution with a slight positive skew.*


*The histogram confirms this, as most ratings are clustered near the
mean with a small rightward tail.*

*Conclusion: The nearly symmetrical distribution suggests balanced user
satisfaction.*

*This indicates that the app performs moderately well in meeting user
expectations.*

...

```
plt.figure(figsize=(4,4))
sns.histplot(data = df["Rating"], bins = 5, kde = True, edgecolor =
"black")
plt.title("Distribution of App Ratings", fontsize=14)
plt.show()
skew = df["Rating"].skew()
print("Skewness : ",skew)
```



Skewness : 0.10182054838079216

#Hypothesis testing

```
instagram_ratings = df[df["App_Name"] == "Instagram"]["Rating"]  
whatsapp_ratings = df[df["App_Name"] == "WhatsApp"]["Rating"]  
instagram_ratings.describe()
```

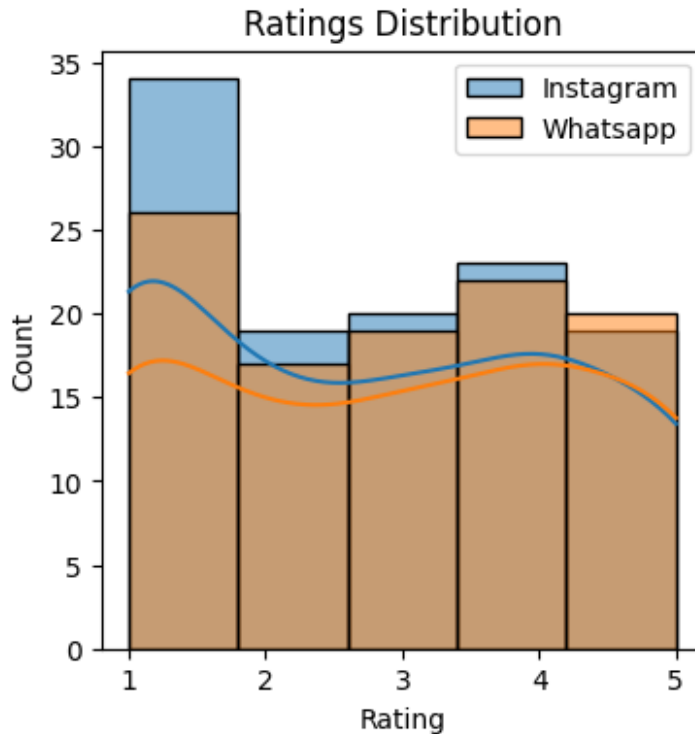
```
count      115.000000  
mean        2.773913  
std         1.475299  
min         1.000000  
25%         1.000000  
50%         3.000000  
75%         4.000000  
max         5.000000  
Name: Rating, dtype: float64
```

```
whatsapp_ratings.describe()
```

```
count      104.000000  
mean        2.932692  
std         1.469855  
min         1.000000  
25%         1.750000  
50%         3.000000  
75%         4.000000  
max         5.000000  
Name: Rating, dtype: float64
```

plot for ratings distribution

```
plt.figure(figsize = (4,4))  
sns.histplot(data = instagram_ratings, bins= 5, kde= True, label =  
"Instagram")  
sns.histplot(data = whatsapp_ratings, bins= 5, kde= True, label =  
"Whatsapp")  
plt.title("Ratings Distribution")  
plt.legend()  
plt.show()
```



```

stat, p = levene(instagram_ratings, whatsapp_ratings)
print("Levene's Test P-Value:", p)
print("-----")
t_stat, p_value = ttest_ind(instagram_ratings, whatsapp_ratings,
                             alternative='greater')

print("T-Statistic:", t_stat)
print("P-Value:", p_value)

Levene's Test P-Value: 0.7874164357527333
-----
T-Statistic: -0.79674231444911
P-Value: 0.786764229580496

#Performing CLT

#Parameters
sample_size = 30
num_samples = 1000

# Population data = Ratings column
population = df["Rating"]

# Step -1 ; Taking random samples and calculating their means
sample_means = []
for i in range (num_samples):
    sample = np.random.choice(population,size = sample_size, replace =

```

```

True)
    sample_means.append(np.mean(sample))


#step -2 ; Creating sampling distribution
plt.figure(figsize = (4,4))
sns.histplot(data = sample_means,bins = 30, edgecolor = "black", kde =
True)
plt.title("Sampling distribution of the Mean (sample size = 30)",
fontsize = 12)
plt.show()

#Step - 3 ; Comparing Population mean with Sampling mean
population_mean = np.mean(population)
sampling_mean = np.mean(sample_means)
sampling_std = np.std(sample_means)

print("Population Mean:", population_mean)
print("Sampling Mean:", sampling_mean)
print("Standard Error:", sampling_std)

```

Sampling distribution of the Mean (sample size = 30)



```

Population Mean: 2.869
Sampling Mean: 2.8728000000000002
Standard Error: 0.2718294890388295

```